

Foreword

I used the c programming language to complete this task. I represented the system of spins as a struct, defined below in one dimension.

```
typedef struct ising_t
{
    float epsilon;
    float magnetic_field;
    float epsilon;
    int length;
    int *ensemble;
} ising_t;
```

Scaling this to two dimensions was more difficult because the `*ensemble` had to be transformed into `**ensemble`. It is possible to implement this in a single struct using union `ensemble {int *1d, int **2d};`, but in the code available on my github I have used two separate structs resulting in code duplication for many of the methods. I made this decision because I had not used c prior to this project and was unaware of the union keyword and its uses. I am describing this because you may notice discrepancies in the code snippets that I have include where `int *(*)ensemble = system -> ensemble;` is invoked to access the array of spins.

Question 1 a)

I selected three different temperatures, $1.0, 2.0$ and $3.0\epsilon/k$, and ran the metropolis algorithm for $1000N$ steps. At each of these temperatures the system was initialised in a random state and allowed to equilibrate.




Figure 1: Each vertical pair of lines represents the spin state. The top one is the random initial state and the bottom one is the equilibrated final state.

We noticed that at lower temperatures the spin chunks

were much larger than at higher temperatures. This is particularly pronounced between the $T = 3.0\epsilon/k$ and $T = 1.0\epsilon/k$ plots in figure 1. To evolve the system we used the version of the metropolis algorithm shown below,

```
/*
 * metropolis_step
 * -----
 * Evolve the system by attempting to
 * flip a spin. The step is weighted
 * by the Boltzmann factor.
 *
 * parameters
 * -----
 * ising_t *system: A struct
 * encapsulating the information
 * related to the system.
 */
void metropolis_step(ising_t *system)
{
    float temp = system -> temperature;
    int *ensemble = system -> ensemble;
    int length = system -> length;
    int spin = random_index(length);
    int change = 2*ensemble[spin]*(
        ensemble[modulo(spin+1,length)] +
        ensemble[modulo(spin-1,length)]);

    if ((change < 0)
        || (exp(-change/temp)>randn()))
    {
        ensemble[spin] *= -1;
    }
}
```

where, `rand` generates a random number in the range $[0,1]$, and `modulo` is modified to produce positive input on negative numbers like the python implementation. This is not the native c implementation. I used the `||` short circuit operator so that the second comparison was not evaluated on every call to the function.

Question 1 b)

I found that it was worth considering what the *basic unit* of the Ising model was. In the absence of an external magnetic field the energy is a function of the pairs. I start by considering the partition function of an individual pair. This is a two level system; either the pair are aligned or they are anti-aligned with the corresponding energies.

$$\begin{aligned} Z_i &= \sum_{s_i=\pm 1} \exp\left(-\frac{\epsilon s_i s_{i+1}}{\tau}\right) \\ &= \exp\left(-\frac{\epsilon}{\tau}\right) + \exp\left(\frac{\epsilon}{\tau}\right) \\ &= 2 \cosh\left(\frac{\epsilon}{\tau}\right). \end{aligned} \quad (1)$$

Similarly to the para-magnetic case we can multiply the system partition functions of single constituents together

to get the partition function of the entire system. However, the condition to do this was that the constituents were independent, but the Ising model contains interactions. In the case of the Ising model the constituents that are independent are the pairs, not the individual spins. You may think then that we only consider $N/2$ unique pairs but this is not the case. In a chain each spin is counted in two pairs so the power is still N .

A small detail that I skipped was what happens at the boundary. The two spins on the end of the chains are not (necessarily) counted twice. In the limit of a very large chain of spins we can see that the boundary affect will not matter however, we got about this nuance in a much more interesting way by considering cyclic boundary conditions. That is to say that the spin on the far end of the chain is a neighbour to the spin at the start of the chain and vice versa.

Given the partition function $Z = (2 \cosh(\varepsilon/\tau))^N$, we calculated the internal energy using,

$$\begin{aligned}
 U &= \tau^2 \partial_\tau \ln(Z) \\
 &= \tau^2 \partial_\tau \ln \left(2 \cosh \left(\frac{\varepsilon}{\tau} \right)^N \right) \\
 &= N \tau^2 \partial_\tau \ln \left(2 \cosh \left(\frac{\varepsilon}{\tau} \right) \right) \\
 &= N \tau^2 \partial_\tau \left(2 \cosh \left(\frac{\varepsilon}{\tau} \right) \right) \frac{1}{2 \cosh \left(\frac{\varepsilon}{\tau} \right)} \\
 &= N \tau^2 \partial_\tau \left(\frac{\varepsilon}{\tau} \right) \frac{\sinh \left(\frac{\varepsilon}{\tau} \right)}{\cosh \left(\frac{\varepsilon}{\tau} \right)} \\
 &= -\varepsilon N \tanh \left(\frac{\varepsilon}{\tau} \right).
 \end{aligned} \tag{2}$$

We calculated the free energy of the system using,

$$\begin{aligned}
 F &= -\tau \ln Z \\
 &= -\tau \ln \left(\left(2 \cosh \left(\frac{\varepsilon}{\tau} \right) \right)^N \right) \\
 &= -N \tau \ln \left(2 \cosh \left(\frac{\varepsilon}{\tau} \right) \right) \\
 &= -N \tau \ln \left(\exp \left(\frac{\varepsilon}{\tau} \right) + \exp \left(-\frac{\varepsilon}{\tau} \right) \right) \\
 &= -N \tau \ln \left(\exp \left(\frac{\varepsilon}{\tau} \right) \left(1 + \exp \left(-2 \frac{\varepsilon}{\tau} \right) \right) \right) \\
 &= -N \tau \ln \left(\exp \left(\frac{\varepsilon}{\tau} \right) \right) - N \tau \ln \left(1 + \exp \left(-2 \frac{\varepsilon}{\tau} \right) \right) \\
 &= -N \varepsilon - N \tau \ln \left(1 + \exp \left(-2 \frac{\varepsilon}{\tau} \right) \right).
 \end{aligned} \tag{3}$$

The entropy followed from the combination of Equation 13 and Equation 12 using Equation 14,

$$\tau \sigma = F - U \tag{4}$$

$$\begin{aligned}
 &= -N \varepsilon \tanh \left(\frac{\varepsilon}{\tau} \right) + N \varepsilon + N \tau \ln \left(1 + \exp \left(-2 \frac{\varepsilon}{\tau} \right) \right) \\
 \sigma &= \frac{\varepsilon}{\tau} \left(1 - \tanh \left(\frac{\varepsilon}{\tau} \right) \right) + \ln \left(1 + \exp \left(-2 \frac{\varepsilon}{\tau} \right) \right).
 \end{aligned} \tag{5}$$

Finally, we determined the specific heat using Equation 13 and Equation 14,

$$C = \partial_\tau U \tag{6}$$

$$\begin{aligned}
 &= \partial_\tau \left(-N \varepsilon \tanh \left(\frac{\varepsilon}{\tau} \right) \right) \\
 &= -N \varepsilon \partial_\tau \left(\frac{\varepsilon}{\tau} \right) \frac{1}{\cosh^2 \left(\frac{\varepsilon}{\tau} \right)} \\
 &= \frac{N \varepsilon^2}{\tau^2 \cosh^2 \left(\frac{\varepsilon}{\tau} \right)}.
 \end{aligned} \tag{7}$$

Question 1 c)




Figure 2: The top left is the energy and the top right is the entropy. The bottom left is the free energy and the bottom right is the heat capacity.

I started by simulating a one dimensional Ising model with no external magnetic field, which I compared to the analytic expressions derived above. I used periodic boundary conditions and chose to implement my models using a lattice size of one-hundred spins. I chose to use one-hundred spins because it evaluated fast on my device and was large enough to be interesting.

Starting with our one dimensional model we equilibrated the system for multiple different temperatures and settled on using $1000N$ as the length of the loop. This was likely too many but I found that for low temperatures when the probability of a flip becomes small, a larger number of steps was required.

I chose to sample the temperatures over the range $0.0 - 4.0\epsilon/k$ incrementing by $0.2\epsilon/k$. I initialised the system only once at the highest temperature that we sampled, $3.8\epsilon/k$. I equilibrated the system at this temperature by evolving it for $1000N$ and then started to cool the system taking measurements at each new system.

The alternative model was to randomly initialise the system at every temperature. This would require approximately twice the number of steps since the system would have to be equilibrated at every temperature. I realize that the cooling method has a side affect of leading to "overflow". By "overflow" I mean that the first few measurements of each temperature are slightly out of equilibrium at the higher temperature.

The entire *measured* cooling process was performed by the program 1000 times. At each temperature in the *measured cooling* loop the energy and entropy were measured by taking the average of all $1000N$ iterations. The heat capacity was also measured by taking the variance of the energy and applying,

$$C_v = \frac{\text{var}(\epsilon)}{\tau^2} \quad (10)$$

The energy and entropy were re-averaged over the outer *fixed size* loop. I chose to re-average then on the outer loop because I was certain that the trials were statistically independent. This does not matter so much for the mean value estimate since the $1000N$ inner loop allows the system to explore the equilibrium space but I think that it does matter for the error estimate.

I believe that it matters for the error estimate because the state of the system is highly correlated to the past state of the system within some *correlation length*. This *correlation length* is roughly the same amount of time that the system requires to explore the equilibrium state, or $1000N$.

By running multiple simulations for the *correlation length* and averaging these results I have guaranteed robust results. I estimated the error using the *standard error* of the independent *correlation length* simulations. This means that the error presented in figure 2 is given by equation 17.

$$\Delta\hat{\beta} = \sqrt{\frac{\text{var}(\beta)}{N}}, \quad (11)$$

where, $\hat{\beta}$ is the parameter estimate, β is a vector of measurements and N is the number of measurements.

To calculate the energy of the system I used the following algorithm,

```
/*
 * energy_ising_t
 * -----
 * Calculate the energy of the system.
 *
 * parameters
 * -----
```

```
 * ising_t *system: A struct
 *   encapsulating the information
 *   related to the system.
 *
 * returns
 * -----
 * float energy: The energy of the
 *   system in Joules.
 */
float energy_ising_t(ising_t *system)
{
    int length = system->length;
    int *ensemble = system->ensemble;
    float energy = 0.;

    for(int spin=0; spin<length; spin++)
    {
        energy -= ensemble[spin] *
            ensemble[modulo(spin+1,length)];
    }

    return energy;
}
```

You may notice that I am only counting the righthand neighbour of each spin. I chose this method because it is more optimal. If I was to count each neighbour then every pair would be counted twice and we would have to divide the final result by 2.. By only counting one of the neighbours I have halved the number of computations.

A pair of spins is the *base unit* of the ising model so to calculate the entropy I counted the number of aligned pairs and then used the *chose* function to calculate the multiplicity. However, it was not quite this simple since $100!$ is $\sim 10^{157}$, which overflows an integer in the programs memory.

To fix this problem I used the stirling approximation to compute the entropy directly. This implies that the entropy should be less accurate at lower temperatures where the entropy is low and the Stirling approximation diverges from the actual entropy. The code that I used to calculate the entropy was,

```
/*
 * entropy_ising_t
 * -----
 * Calculate the entropy of a
 * configuration.
 *
 * parameters
 * -----
 * ising_t *system: A struct
 *   encapsulating the information
 *   related to the system.
 *
 * returns
 * -----
 * float entropy: The entropy of the
 *   system in natural units.
```

```

*/
float entropy_ising_t(ising_t *system)
{
    int length = system->length;
    int *ensemble = system->ensemble;
    int up = 0;

    for (int spin=0; spin<length; spin++)
    {
        up += ensemble[spin] ==
            ensemble[modulo(spin+1,length)];
    }

    int down = length - up;

    float entropy = length * log(length) -
        up * log(up) - down * log(down);

    return entropy;
}

```

Question 1 d)

Consider the energy depicted in figure 2. We see that the energy is a decreasing function of the temperature. This implies that the spins tend to align as the temperature decreases. This makes sense because the Boltzmann factor for the lower state becomes more favoured. Similarly the entropy is a increasing function of the temperature. As the spins tend to align at lower temperatures the number of ways to arrange the state becomes smaller, decreasing the entropy.

The heat capacity, also shown in figure 2 is interesting because it increases rapidly before asymptotically decreasing. If it was discontinuous then we would expect a phase transition however it is continuous. so we know that there is no phase transition. We know it is continuous because of equation 16.

Question 1 e)

Figure 3 showed me that the magnetisation was not zero as it was predicted to be at $0\epsilon/k$. I am not concerned because the distribution of samples, which roughly corresponds to a time average is symmetrically distributed around zero implying that the time average is identically zero. I would expect some variation of the results around zero as spins randomly flipped.

Since magnetisation sets in at a lower temperature for the larger sample I concluded that there was no phase transition. I decided this because the phase transition should occur at the same temperature for all lattice sizes. Taking the limit of a decreasing function presumably reduces to zero implying that there is no spontaneous magnetisation for the one-dimensional Ising model. This means that there is no phase transition.




Figure 3: Histograms of the magnetisation for two ising systems of sizes $N = 100$ and $N = 500$. The temperatures are labelled in the titles and the y -axis is a unitless count.

Question 2 a)

For $\tau = 1.0\epsilon/k$ the random lattice very quickly relaxed into large *domains* of up and down spins. Over time the boundaries of the *domains* moved around the lattice coherently. I am/will using/use *coherent motion* to refer to a blob transporting across the network. For example, a group of down spins may start in the top right corner and remaining a single blob move to the middle of the system.

I noticed that given enough time the $\tau = 1.0\epsilon/k$ lattice relaxed until it was entirely one color. This is not shown in figure ?? simply because the lattice was not allowed to evolve for enough time. I used the same evolution time for the two dimensional Ising model that I used for the one dimensional Ising model, that is $1000N^2$ where I will use N to represent the side length of the system.

This is important because at lower temperatures the system took longer to equilibrate. As I saw for the one-dimensional case the *coherent domains* have a *meta-stable* lifetime before decaying. The shape of this domain influences the lifetime. The size seems to be a direct predictor of lifetime with the largest domain nearly always coming out on top. Another important feature that we noticed was that the coherent domains became particularly *meta-stable* if they formed rings.

The ring structures were *meta-stable* because of the cyclic boundary conditions. In a ring structure there are only two edges that are exposed to the other *meta-stable* state meaning that the anti-aligned interactions at the boundary are minimised. Additionally since the interaction depth of the Ising model is only nearest neighbour the *meta-stable* ring could exit in a band of just five spins. *The spins on the boundary have no idea how large their respective domains are.*

A final point to address for the $\tau = 1.0\epsilon/k$ model is the *time speckling*. As the system is allowed to evolve in equilibrium spins at random locations will randomly flip. This

is a byproduct of the metropolis algorithm but it is important because the *speckles* almost always remain a single spin. You will see later that this changes as the temperature is increased.

Heating the lattice to $\tau = 2.0\epsilon/k$ we see much of the same behaviour. $\tau = 2.0\epsilon/k$ is still less than the critical temperature, so all of the same evolutionary events occur. *Meta-stable* domains form, move and die. Even the more dangerous *meta-stable* rings can form at this temperature. I would like to say that the lifetime of the *domains* has declined but I have no evidence.

The key difference between the $\tau = 2.0\epsilon/k$ and the $\tau = 1.0\epsilon/k$ models is the *speckling*. At the higher temperatures the speckles occur more often and can result in very local short lived domains. This is why I postulated that the *domains* are shorter lived. The system is more active meaning that the spins on the boundary flip more readily. Moreover, the rapid speckling assists in the destruction of the *domain* when it occurs on the boundary.

Something worthy of note before we move above the condensation temperature is that *domains* die at their corners. This is because the corners represent an interface where a spin has more neighbours of one color than the other. This is another factor that prolongs the life of the *meta-stable* states, since they have no corners and minimised contact.

The behaviour was markedly different for the $\tau = 3.0\epsilon/k$ model, because it is above the critical temperature. For starters the activity of the system has increased significantly. This means, among other things that the relaxation into equilibrium is almost instantaneous compared with the lower temperature models. In addition *domains* do not form at this temperature.

I have been using *domain* to refer to a structure that is *macroscopic* and *semi-permanent*. *Zonation* does occur at $\tau = 3.0\epsilon/k$ but it does not meet these criteria. By this I mean that the zones are small ~ 10 spins maximum and temporary. Rather than considering the rapid zonation to be the evolution of the *domain* behaviour I consider it the evolution of the *speckling* behaviour.

At $\tau = 3.0\epsilon/k$ we also lose the *coherent* end state. At both of the other temperatures the final state was a completely magnetised block or a *meta-stable* ring. However, at $\tau = 3.0\epsilon/k$ there is no such state and the *zonation* appears to randomly move about the grid. At higher temperatures still the *zonation* is completely destroyed and not even microscopic domains are able to form.

Question 2 b)

Smoothly varying the temperature using I found that the *speckling* and *zonation* seems to dominate above $\sim 2.3\epsilon/k$ while below this *domains* start to form and given enough time the entire crystal becomes magnetised. The transition is smooth with respect to the *speckles*, which at very




Figure 4: Snapshots of the Ising system at multiple temperatures given a random initial starting position (top). You will notice that $\tau = 1.0\epsilon/k$ and $\tau = 2.0\epsilon/k$ have similar net magnetisation, but that $\tau = 2.0\epsilon/k$ has stronger *speckling* than $\tau = 1.0\epsilon/k$. You will notice that there is not net magnetisation for $\tau = 3.0\epsilon/k$ and that the *zonation* I described above is occurring.

low temperatures are very small and isolated with out a lifetime.

As I increased the temperature the *speckles* increases in size and in lifetime. A *speckle* is differentiated from a *zone* rather arbitrarily, the two main differences being transport and lifetime. The *zones* have a lifetime of long enough that they are able to *coherently* move across the crystal. On the other hand a *speckle* tends to die where it started and very quickly.

As we approach the critical temperature from below the *speckles* increase in size and disrupt the domains more quickly but, the system is still able to fully magnetise. At the critical temperature the *speckles* become so disruptive that the magnetisation cannot occur macroscopically but it still occurs locally in the form of the *zones* which outlive the now very active *speckles*.

Increasing the temperature further still the coherent *zones* disappear (visually) entirely and the speckles become the dominant behaviour. As a final note it is worth noting that these behaviours I have described were visible to me in part because of the settings of my simulation. I found it was best to run them with many steps per frame 5000 \sim 10000 because it allows you to see the macroscopic behaviours.

I acknowledge that although I wanted to use the qt framework to create my own gui I ran out of time. Instead I used the free online applet created by Daniel V. Schroeder for Weber state university to make the higher level observations. The code that I used to evolve the two dimensional system was,

/*

```

* metropolis_step_ising_t
* -----
* Evolve the system according to a
* randomly weighted spin flip that
* compares the probability of the
* two states based on the Boltzmann
* distribution of the two systems.
*
* parameters
* -----
* ising_t *system: The system to
* evolve.
*/
void metropolis_step(ising_t *system)
{
    int length = system->length;
    int **ensemble = system->ensemble;
    float epsilon = system->epsilon;
    float temp = system->temperature;
    float field = system->magnetic_field;

    int row = random_index(length);
    int col = random_index(length);

    int spin = ensemble[row][col];
    int neighbours =
        ensemble[modulo(row+1,length)][col] +
        ensemble[modulo(row-1,length)][col] +
        ensemble[row][modulo(col+1,length)] +
        ensemble[row][modulo(col-1,length)];

    float magnetic_change=-2*spin*field;
    float interaction_change=2*epsilon*
        neighbours*spin;
    float change=magnetic_change+
        interaction_change;

    if ((energy_change < 0) ||
        (exp(-change/temperature) > randn()))
    {
        ensemble[row][col] *= -1;
    }
}

```

Note: The magnetic_field was set to zero in the provided simulation and discussion, while the epsilon was set to one.

Question 2 c)

I noticed that the heat capacity *spike* increased in height as I increased the size of the crystal. This suggests that the infinite Ising model will have an infinitely tall heat capacity *spike* at the critical temperature. I also noticed that at the critical temperature all of the physical parameters changed the fastest.

This makes sense because a phase transition is loosely defined as a sudden discontinuity in the bulk properties of

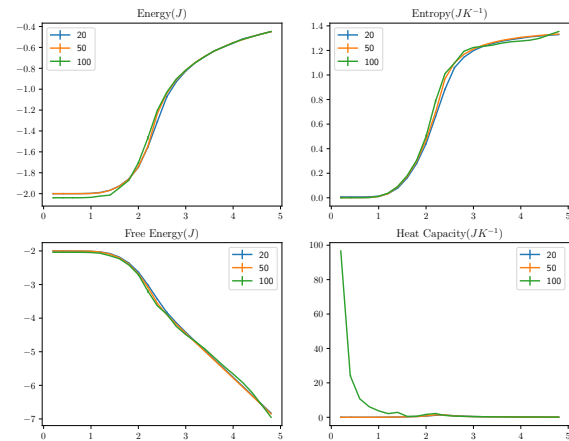


Figure 5:

the sample. Although the entropy and energy are not discontinuous over the critical temperature they do change very suddenly. The discontinuity is in the heat capacity which is the derivative of the energy with respect to temperature. Because a derivative is the discontinuous quantity it implies that the phase transition is second order.

To measure the energy of the two dimensional Ising system I used a very similar piece of code to the one dimensional scenario.

```

/*
* energy_ising_t
* -----
* Calculate the energy of the
* Ising system.
*
* parameters
* -----
* ising_t *system: The system to
* measure.
*/
float energy(ising_t *system)
{
    int length = system->length;
    int **ensemble = system->ensemble;
    float epsilon = system->epsilon;
    float field = system->magnetic_field;
    float magnetic = 0.0;
    float interactions = 0.0;

    for (int row=0; row<length; row++)
    {
        for (int col=0; col<length; col++)
        {
            float neighbours =
                ensemble[modulo(row+1,length)][col] +
                ensemble[modulo(row-1,length)][col] +
                ensemble[row][modulo(col+1,length)] +

```



```

    ensemble[row][modulo(col-1,length)];

    magnetic -= ensemble[row][col] *
        field;
    interactions -= neighbours *
        epsilon * ensemble[row][col];
}
}

return interactions / 2. + magnetic;
}

```

For the entropy the calculation was also very similar.

```

/*
 * entropy_ising_t
 * -----
 * Calculate the entropy of the Ising
 * system.
 *
 * parameters
 * -----
 * ising_t *system: The system to
 * measure.
 */
float entropy_ising_t(ising_t *system)
{
    int len = system->length;
    int **ensemble = system->ensemble;
    int up = 0;

    for (int row=0; row<len; row++)
    {
        for (int col=0; col<len; col++)
        {
            up += ensemble[row][col] ==
                ensemble[modulo(row+1,len)][col];
            up += ensemble[row][col] ==
                ensemble[row][modulo(col+1,len)];
        }
    }

    int total = 2 * len * len;
    int down = total - up;

    return total * log(total) -
        up * log(up) - down * log(down);
}

```

Question 2 d)

I noticed that the magnetisation began to diverge at the same temperature irrespective of the size of the system. This is what I would expect from a phase transition because it should not depend on the size of the sample. To give a real world example iron melts at the same temperature regardless of the size of the block. We might more quickly melt a smaller block but that is just because there is less of it and less self insulation.




Figure 6: The magnetisation of the Ising model as the temperature and number of spins is varied.

As N increases the critical temperature will stay the same, but the steepness of the transition appears to radically increase. I was surprised to see this relationship and I believe it may be because we are at the very low number end of the spectrum. In this regime if one or two spins flip it becomes visible on the graph. What I expect to see is that the lines smoothly converge down to a step function at the critical temperature.

Question 2 e)

I initialised the system randomly. This corresponds to the equilibrium of a system that is at a high temperature. Therefore, I first cooled the system to reach an equilibrium that I would expect for a low temperature state. This was done slowly letting the system equilibrate at each new temperature. This technique avoids the *meta-stable* rings from forming.

The cooled system represents the leftmost figure in 7. We can see that the magnetisation is complete and in this case there are not even any speckles. This is suspiciously good but I cannot find the error in my code so I will run with it. At the highest temperature we see that the state appears to be almost entirely random. Cooling it back down I managed to get a *meta-stable* ring state.

In all the heating and cooling operations I raised the temperature slowly and re-equilibrated the system at each new temperature. I am annoyed by the ring state because this gradual cooling is supposed to prevent such states for the most part. Again this state is markedly non-speckled and I think this is because the temperature is very close to $0.0\epsilon/k$. We also see that the edges of the ring are pixel flat.

I justify this because the corners are where the *domains* die as I discussed above. Hence any incursion that yellow makes into the purple band is unstable because the in-




Figure 7: Heating and cooling an Ising system to change the magnetisation. The leftmost figure is the initial state at the lowest temperature, the middle figure is the highest temperature state and the rightmost figure is the final final state after the system has again been cooled.

curing yellow are more likely to have purple neighbours than yellow neighbours. The same can be said for the purple incursions into yellow territory. Therefore, by far the most stable configuration is the flat edged band which we observe.

The final thing I want to discuss is the change in the magnetisation of the system between the first cool state and the final cool state. We can see that the injection of thermal energy into the system has changed the magnetisation from purple to predominantly yellow ignoring the *meta-stable* ring.

The Currie temperature is defined as the temperature at which the magnetic properties of a metal change dramatically. Since this is the behaviour that we are observing at the critical temperature it makes sense that the critical temperature is analogous to the Curie temperature. In the briefing we were given that the critical temperature for the two dimensional Ising model was $T_c \approx 2.27\epsilon/k$. The Curie temperature of iron is $T_c = 1043K$, hence $\epsilon = 1043K * k/2.27 = 0.0396eV$.

Introduction

The study of magnetic materials is an area of academic and industrial interest (**Annon, dd/mm/yyyy**). For example, magnetic technologies are important in the on-going development of quantum computers, superconducting circuits and other examples in electronics (**Annon, dd/mm/yyyy**). At a fundamental level magnetisation is a well understood phenomenon, yet it is difficult to theoretically model. One simple model of magnetic materials is the Ising model.

The Ising model is the simplest model of a ferro-magnet (**Annon, dd/mm/yyyy**). Despite the simplicity of the Ising model it displays rich physical behaviour and has analytic solutions in one and two dimensions (**Annon, dd/mm/yyyy**). The Ising model is the simplest model to account for inter-molecular interactions and contain a phase transition. This makes it an excellent medium for studying magnetic phenomenon (**Annon, dd/mm/yyyy**).

By modifying the basic Ising model we can simulate many phenomenon including glasses (**Annon, dd/mm/yyyy**). The Ising model has broader significance and can be used to construct very simple neural networks called Boltzmann machines (**Annon, dd/mm/yyyy**). We tested one and two dimension Ising models and confirmed that they matched theoretical predictions.

Theory

Materials have internal interactions. As physicists we like to ignore these where possible but often these approximations limit the accuracies of our models (**Annon, dd/mm/yyyy**). Magnetic phenomenon are no different. To understand how spins interact in a magnet it helps to first construct the simplest possible model without interactions; a para-magnet.

Consider our magnet as a one-dimensional chain of atomic spins. For the moment ignore any external magnetic field and just consider the spins in isolation. Now let's limit the spins to be fixed up or down along one axis. If there are no interactions between the spins the energy is fixed. If we add an external magnetic field then we would expect the ensemble to develop a net magnetisation.

If the system has thermal energy we would expect some of the spins to align themselves anti-parallel to the magnetic field. We can see this affect by considering the partition function for a single spin in the ensemble. If the spin is aligned with the magnetic field then the energy is $-sB$, where s is the unit of magnetisation carried by the single spin and B is the strength of the external magnetic field. If the spin is anti-aligned with the field then the energy is sB .

This is a simple two level system and the partition function

is given by,

$$\begin{aligned} Z &= \sum_{s=\pm 1} \exp\left(-\frac{sB}{\tau}\right) \\ &= \exp\left(-\frac{sB}{\tau}\right) + \exp\left(\frac{sB}{\tau}\right) \\ &= 2 \cosh\left(\frac{sB}{\tau}\right), \end{aligned} \quad (12)$$

where, $\tau = kT$ is the temperature in units of energy. The probability of the spins being anti-aligned with the field is therefore,

$$P = \frac{\exp\left(-\frac{sB}{\tau}\right)}{2 \cosh\left(\frac{sB}{\tau}\right)}. \quad (13)$$

Hence, as the temperature increase we expect the number of anti-aligned spins to increase and as we increase the magnetic field we expect the number of anti-aligned spins to decrease.

Since each of the spins in a para-magnetic system is independent the partition function of an ensemble of N spins is just the product of N partition functions for the single spin case. However, since the spins are indistinguishable we must also divide by a Gibbs correction factor of $N!$. The probability of finding a particular state however, is a case that is worth studying, since it indicates a divergence between the Ising model of a ferro-magnet and a para-magnet in a magnetic field. First we need to define our state.

The energy of the system, and any other physical parameters, only depend on the number of spins that are aligned with the magnetic field and not specifically which spins are aligned with the field. Naively we might expect that the probability of having N_{\uparrow} spins aligned with the field would be,

$$P(N_{\uparrow}) = \frac{\exp\left(-\frac{sN_{\uparrow}B}{\tau}\right) \exp\left(\frac{s(N-N_{\uparrow})B}{\tau}\right)}{\cosh\left(\frac{sB}{\tau}\right)^N}. \quad (14)$$

However, equation 14 has failed to account for the multiple micro-states that occupy this macro-state. We can account for this by multiplying by the multiplicity, which can be found using the chose function,

$$P(N_{\uparrow}) = \frac{N!}{N_{\uparrow}!(N-N_{\uparrow})!} \frac{\exp\left(-\frac{sN_{\uparrow}B}{\tau}\right) \exp\left(\frac{s(N-N_{\uparrow})B}{\tau}\right)}{\cosh\left(\frac{sB}{\tau}\right)^N}. \quad (15)$$

Equation 15 is the correct expression for the probability.

It is informative to calculate the internal energy and free energy of the system. Starting with the internal energy,

$$\begin{aligned} U &= \tau^2 \partial_{\tau} \ln Z \\ &= \tau^2 \partial_{\tau} \ln \left(2^N \cosh^N \left(\frac{sB}{\tau} \right) \right) \\ &= -NsB \tanh \left(\frac{sB}{\tau} \right). \end{aligned} \quad (16)$$

We can also calculate the free energy, but further calculations result in tedious analytical expressions so we have omitted them.

$$\begin{aligned} F &= -\tau \ln Z \\ &= -\tau \ln \left(2^N \cosh^N \left(\frac{sB}{\tau} \right) \right) \\ &= -NsB - N\tau \ln \left(1 + \exp \left(-\frac{2sB}{\tau} \right) \right). \end{aligned} \quad (17)$$

As we will see when we analyse the Ising model without an external field these results are general of any two level system. Using equation 16 we can calculate the magnetisation as a function of the magnetic field and temperature,

$$\begin{aligned} U &= mB = -NsB \tanh \left(\frac{sB}{\tau} \right) \\ m &= -Ns \tanh \left(\frac{sB}{\tau} \right). \end{aligned} \quad (18)$$

Therefore, the net magnetisation system will decrease with temperature and increase with the magnetic field, much as we would expect.

Para-magnets are a useful toy model but from our experience with natural and manufactured magnets we know that it is possible to construct systems that are magnetic without external fields. The one-dimensional Ising model is a simple model of such systems. The Ising model is a natural extension of the paramagnetic model that we discussed, and operates on the same spin lattice.

The Ising model differs because it adds very simple interactions between neighbouring spins. This interaction favours pairs that are aligned by reducing the energy of this scenario. Representing up spins as $+1$ and down spins as -1 we can represent this mutual interaction as $\Delta\epsilon = \epsilon s_i s_{i+1}$, where $\Delta\epsilon$ is the energy contribution of the interaction, ϵ is a scaling factor that represents the strength of the interaction and s_i is the i^{th} spin in the chain.

What happens if we place the Ising model into an external magnetic field. Again we can break it down by considering a single pair in the chain as our constituent object. There are three energies that it is possible for this pair to have; parallel and aligned with the magnetic field, parallel and anti-aligned with the magnetic field and anti-parallel. However, the final state has a multiplicity of two since either of the spins could be aligned with the field.

It is possible to compute the partition function for this the single pair and hence also the entire system.

$$\begin{aligned} Z_1 &= \sum_s \exp \left(-\frac{\epsilon_s}{\tau} \right) \\ &= \exp \left(\frac{-\epsilon - 2B}{\tau} \right) + 2 \exp \left(\frac{\epsilon}{\tau} \right) + \exp \left(\frac{-\epsilon + 2B}{\tau} \right) \\ &= 2 \exp \left(\frac{-\epsilon}{\tau} \right) \cosh \left(\frac{2B}{\tau} \right) + 2 \exp \left(\frac{\epsilon}{\tau} \right). \end{aligned} \quad (19)$$

From here we can compute all the physical properties of the system. Since each pair is independent the partition function is simply the product of N partition function by the same reasoning as in the case above when there was no magnetic field. We have not shown the calculation of the energy etc., because the expressions are complex and uninformative.

Another interesting effect that can be explored using the Ising model is anti-ferro-magnetism. This phenomenon was only recently discovered in nature (**Annon, dd/mm/yyyy**) and refers to an interaction between neighbouring spins that causes them to have lower energy when they are aligned anti-parallel rather than parallel. We do not need to cover any new equations in this case as an anti-ferro-magnet can be explored by letting ϵ become negative.

We have spent a lot of time discussing the one-dimensional scenario but real systems are typically higher dimensional. There is an analytical solution to the two-dimensional Ising model (**Annon, dd/mm/yyyy**). This solution is a tour de force and has comparatively little practical use due to its complexity. Multiple approximation methods have been developed for dealing with the two-dimensional case, most notably the mean field approximation.

The mean field approximation treats a group of neighbours as a single spin, parametrised by the mean. In this way we recover the two level system and arrive at a two level system that is very similar to what we have already covered for the para-magnet and the Ising model when there is no external magnetic field. Ultimately the mean field approximation is an approximation and its predictions are not always correct.

Another aspect of higher dimensions that it is worth discussing is what counts as a neighbour. For example, in two-dimensions we could connect the spins together so that each spin is equally far from six other spins. This triangular Ising model will have markedly different behaviour than a square grid of spins (**Annon, dd/mm/yyyy**). For our analysis we have considered a square grid of spins since it is simpler to simulate.

Without going into the complex analytical solutions we can still make useful qualitative guesses about the behaviour of the Ising model in higher dimensions based on its behaviour in lower dimensions. For example, we expect the spins will tend to align at lower temperatures and tend to disorder at higher temperatures. For the anti-ferro-magnetic case we expect the spins to become anti-aligned at low temperatures and tend to disorder at higher temperatures.

It is worth noting that in the presence of a magnetic field the qualitative behaviour of the anti-ferro-magnetic and ferro-magnetic Ising models becomes markedly different. The ferro-magnetic model results in a positive feedback loop as the magnetic field coerces spins to align with the field they also want to align with each other. On the other hand the anti-ferro-magnet exerts a dampening effect for

the opposite reason.

Method

```
function calc_energy
energy = 0
for spin in 0:length
energy += ensemble[spin] * ensemble[spin + 1]
return energy
```

Similarly for the entropy we counted all of the aligned pairs. This works because as discussed in the theory the "base unit" of the Ising model is a pair of spins not an individual spin. As noted in the Theory there are N pairs of spins. Moreover, we used Stirling's approximation in the logarithmic form, because we found that the program could not calculate numbers of the size $100!$. This will have had negligible effects at higher temperatures where the system tends toward disorder, but at lower temperatures the approximation becomes less accurate. We were not too concerned with the loss of accuracy since the entropy tends to zero at low temperatures.

```
function calc_entropy
up = 0
for spin in 0:length
up += ensemble[spin] == ensemble[spin + 1]
down = length - up
entropy = length * log(length) - up * log(up) - down * log(down)
return entropy
```

The free energy was calculated using its definition, $F = U - \tau\sigma$, where $\tau = kT$ and σ is the entropy. The heat capacity was calculated using the thermodynamic identity,

$$C_V = \frac{\text{var}(U)\tau^2}{k^2 T^2} \quad (20)$$

A large part of the computational expense came from estimating the uncertainty in the physical parameters. It did not make sense to initialise the system randomly at every temperature and then wait for it to equilibrate before taking measurements. Instead we initialised the system at a high temperature where the random configuration is a good approximation to the equilibrium configuration and equilibrated it.

From this higher temperature we allowed the system to evolve for $1000N$ steps measuring the physical properties at every step. We then cooled the system by a small increment and without re-equilibrating the system evolved it for $1000N$ steps taking measurements every step. In this way we halved the amount of CPU time required by each run but introduced a small error by starting the system at a slightly out of equilibrium state. Given the large number of runs we believe that this error is negligible, although it is visible at lower temperatures as the heat capacity becomes over-estimated.

Another flaw of running the simulation this way was that it serialised the loops (i.e. made the next iteration depend on the state of the previous one). This prevented us from making efficient parallelisation of the inner loop, however, the outer loop was not serial and could be efficiently parallelised. After each run of $1000N$ the average was taken for each physical parameter.

Please note that we did not use the standard error of the $1000N$ trials as the error. It does not make sense to do so because each state is deterministically dependent on the previous one and the sample gridding is much too fine. Instead we repeated the entire process a fixed number of times 100 for the one-dimensional case and used the standard error of the means from these 100 trials as the estimate of our uncertainty. Note that we defined the standard error as $\sqrt{\text{Var}N}$.

We noticed that the infinite one-dimensional Ising model is predicted to have no net magnetisation at $0K$. To test this hypothesis, we created histograms of the magnetisation at $\tau = 0.5J, 1.0J$ and $2.0J$ for $N = 100$ and $N = 500$. We expected the histograms to be narrower as the number of spins in the system was increased.

Once we were satisfied with the one-dimensional Ising model we repeated a similar analysis for the two-dimensional model. First we tested the time required for equilibration by initialising the model and running it for $1000N^2$ where N is the width of the grid.

Satisfied with the equilibration time we employed the same techniques described for the one-dimensional case to the two-dimensional case to measure the physical parameters. In short, we incrementally cooled the system measuring every iteration and recording the mean. We repeated this process a fixed number of times and used the standard error as our uncertainty estimate. The energy calculation is clearly modified, so we have included the relevant pseudocode below:

Results

We noticed that the net magnetisation set in at lower temperatures for the larger N . This tells us that there is no phase transition because a phase transition should occur at exactly the same temperature for all lattice sizes. Moreover, it agrees with the theoretical prediction that the Ising model is not magnetised at $0K$ in the infinite case because if we increase the lattice temperature to infinity then the temperature of net magnetisation should decrease below $0K$.