Foreword

I used the `c` programming language to complete this task. I represented the system of spins as a `struct`, defined below in one dimension.

```c
typedef struct ising_t
{
  float epsilon;
  float magnetic_field;
  float epsilon;
  int length;
  int *ensemble;
} ising_t;
```

Scaling this to two dimensions was more difficult because the `*ensemble` had to be transformed into `**ensemble`. It is possible to implement this in a single `struct` using `union ensemble {int *1d, int **2d};`, but in the code available on my `github` I have used two separate `structs` resulting in code duplication for many of the methods. I made this decision because I had not used `c` prior to this project and was unaware of the `union` keyword and its uses. I am describing this because you may notice discrepancies in the code snippets that I have include where `int *(*)ensemble = system -> ensemble;` is invoked to access the array of spins.

Question 1 a)

I selected three different temperatures, $1.0, 2.0$ and $3.0\epsilon/k$, and ran the metropolis algorithm for $1000N$ steps. At each of these temperatures the system was initialised in a random state and allowed to equilibrate.


Figure 1: Each vertical pair of lines represents the spin state. The top one is the random initial state and the bottom one is the equilibrated final state.

We noticed that at lower temperatures the spin chunks were much larger than at higher temperatures. This is particularly pronounced between the $T = 3.0\epsilon/k$ and $T = 1.0\epsilon/k$ plots in figure 1. To evolve the system we used the version of the metropolis algorithm shown below,

```c
/*
 * metropolis_step
 * -----------------------
 * Evolve the system by attempting to
 * flip a spin. The step is weighted
 * by the Boltzmann factor.
 *
 * parameters
 * ----------
 * ising_t *system: A struct
 *   enscapulating the information
 *   related to the system.
 */
void metropolis_step(ising_t *system)
{
  float temp = system -> temperature;
  int *ensemble = system -> ensemble;
  int length = system -> length;
  int spin = random_index(length);
  int change = 2*ensemble[spin]*(
    ensemble[modulo(spin+1,length)] +
    ensemble[modulo(spin-1,length)]);

  if ((change < 0)
    || (exp(-change/temp)>randn()))
  {
    ensemble[spin] *= -1;
  }
}
```

where, `rand` generates a random number in the range $[0, 1]$, and `modulo` is modified to produce positive input on negative numbers like the `python` implmentation. This is not the native `c` implementation. I used the `||` short circuit operator so that the second comparison was not evaluated on every call to the function.

Question 1 b)

I found that it was worth considering what the *basic unit* of the Ising model was. In the absence of an external magnetic field the energy is a function of the pairs. I start by considering the partition function of an individual pair. This is a two level system; either the pair are aligned or they are anti-aligned with the corresponding energies.

$$
\begin{aligned}
Z_i &= \sum_{s_i = \pm 1} \exp\left(-\frac{\varepsilon s_i s_{i+1}}{\tau}\right) \\
&= \exp\left(-\frac{\varepsilon}{\tau}\right) + \exp\left(\frac{\varepsilon}{\tau}\right) \\
&= 2\cosh\left(\frac{\varepsilon}{\tau}\right). \qquad (1)
\end{aligned}
$$

Similarly to the para-magnetic case we can multiply the system partition functions of single constituents together

to get the partition function of the entire system. However, the condition to do this was that the constiutuents were independent, but the Ising model contains interactions. In the case of the Ising model the constituents that are independent are the pairs, not the individual spins. You may think think then that we only consider $N/2$ unique pairs but this is not the case. In a chain each spin is counted in two pairs so the power is still N.

A small detail that I skipped was what happens at the boundary. The two spins on the end of the chains are not (neccessarily) counted twice. In the limit of a very large chain of spins we can see that the boundary affect will not matter however, we got about this nuance in a much more interesting way by considering cyclic boundary conditions. That is to say that the spin on the far end of the chain is a neighbour to the spin at the start of the chain and vice versa.

Given the partition function $Z = (2\cosh(\varepsilon/\tau))^N$, we calculated the internal energy using,

$$U = \tau^2 \partial_\tau \ln(Z) \tag{2}$$
$$= \tau^2 \partial_\tau \ln\left(2\cosh\left(\frac{\varepsilon}{\tau}\right)^N\right)$$
$$= N\tau^2 \partial_\tau \ln\left(2\cosh\left(\frac{\varepsilon}{\tau}\right)\right)$$
$$= N\tau^2 \partial_\tau \left(2\cosh\left(\frac{\varepsilon}{\tau}\right)\right)\frac{1}{2\cosh\left(\frac{\varepsilon}{\tau}\right)}$$
$$= N\tau^2 \partial_\tau \left(\frac{\varepsilon}{\tau}\right)\frac{\sinh\left(\frac{\varepsilon}{\tau}\right)}{\cosh\left(\frac{\varepsilon}{\tau}\right)}$$
$$= -\varepsilon N \tanh\left(\frac{\varepsilon}{\tau}\right). \tag{3}$$

We calculated the free energy of the system using,

$$F = -\tau \ln Z \tag{4}$$
$$= -\tau \ln\left(\left(2\cosh\left(\frac{\varepsilon}{\tau}\right)\right)^N\right)$$
$$= -N\tau \ln\left(2\cosh\left(\frac{\varepsilon}{\tau}\right)\right)$$
$$= -N\tau \ln\left(\exp\left(\frac{\varepsilon}{\tau}\right) + \exp\left(-\frac{\varepsilon}{\tau}\right)\right)$$
$$= -N\tau \ln\left(\exp\left(\frac{\varepsilon}{\tau}\right)\left(1 + \exp\left(-2\frac{\varepsilon}{\tau}\right)\right)\right)$$
$$= -N\tau \ln\left(\exp\left(\frac{\varepsilon}{\tau}\right)\right) - N\tau \ln\left(1 + \exp\left(-2\frac{\varepsilon}{\tau}\right)\right)$$
$$= -N\varepsilon - N\tau \ln\left(1 + \exp\left(-2\frac{\varepsilon}{\tau}\right)\right). \tag{5}$$

The entropy followed from the combination of Equation 3 and Equation 1 using Equation 5,

$$\tau\sigma = F - U \tag{6}$$
$$= -N\varepsilon \tanh\left(\frac{\varepsilon}{\tau}\right) + N\varepsilon + N\tau \ln\left(1 + \exp\left(-2\frac{\varepsilon}{\tau}\right)\right)$$
$$\sigma = \frac{\varepsilon}{\tau}\left(1 - \tanh\left(\frac{\varepsilon}{\tau}\right)\right) + \ln\left(1 + \exp\left(-2\frac{\varepsilon}{\tau}\right)\right). \tag{7}$$

Finally, we determined the specific heat using Equation 3 and Equation 5,

$$C = \partial_\tau U \tag{8}$$
$$= \partial_\tau \left(-N\varepsilon \tanh\left(\frac{\varepsilon}{\tau}\right)\right)$$
$$= -N\varepsilon \partial_\tau \left(\frac{\varepsilon}{\tau}\right)\frac{1}{\cosh^2\left(\frac{\varepsilon}{\tau}\right)}$$
$$= \frac{N\varepsilon^2}{\tau^2 \cosh^2\left(\frac{\varepsilon}{\tau}\right)}. \tag{9}$$

Question 1 c)

Figure 2: The top left is the energy and the top right is the entropy. The bottom left is the free energy and the bottom right is the heat capacity.

I started by simulating a one dimensional Ising model with no external magnetic field, which I compared to the analytic expressions derived above. I used periodic boundary conditions and chose to implement my models using a lattice size of one-hundred spins. I chose to use one-hundred spins because it evaluated fast on my device and was large enough to be interesting.

Starting with our one dimensional model we equilibrated the system for multiple different temperatures and settled on using $1000N$ as the length of the loop. This was likely too many but I found that for low temperatures when the probability of a flip becomes small, a larger number of steps was required.

2

I chose to sample the temperatures over the range $0.0 - 4.0\epsilon/k$ incrementing by $0.2\epsilon/k$. I initialised the system only once at the highest temperature that we sampled, $3.8\epsilon/k$. I equilibrated the system at this temperature by evolving it for $1000N$ and then started to cool the system taking measurements at each new system.

The alternative model was to randomly initialise the system at every temperature. This would require approximately twice the number of steps since the system would have to be equilibrated at every temperature. I realize that the cooling method has a side affect of leading to "overflow". By "overflow" I mean that the first few measurements of each temperature are slightly out of equilibrium at the higher temperature.

The entire *measured* cooling process was performed by the program 1000 times. At each temperature in the *measured cooling* loop the energy and entropy were measured by taking the average of all $1000N$ iterations. The heat capacity was also measured by taking the variance of the energy and applying,

$$C_v = \frac{\text{var}(\epsilon)}{\tau^2} \tag{10}$$

The energy and entropy were re-averaged over the outer *fixed size* loop. I chose to re-average then on the outer loop because I was certain that the trials were statistically independent. This does not matter so much for the mean value estimate since the $1000N$ inner loop allows the system to explore the equilibrium space but I think that it does matter for the error estimate.

I belive that it matters for the error estimate because the state of the system is highly correlated to the past state of the system within some *correlation length*. This *correlation length* is roughly the same amount of time that the system requires to explore the equilibrium state, or $1000N$.

By running multiple simulations for the *correlation length* and averaging these results I have guaranteed robust results. I estimated the error using the *standard error* of the independent *correlation length* simulations. This means that the error presented in figure 2 is given by equation 11.

$$\Delta\hat{\beta} = \sqrt{\frac{\text{var}(\beta)}{N}}, \tag{11}$$

where, $\hat{\beta}$ is the parameter estimate, β is a vector of measurements and N is the number of measurements.

To calculate the energy of the system I used the following algorithm,

```
/*
 * energy_ising_t
 * --------------
 * Calculate the energy of the system.
 *
 * parameters
 * ----------
```

```
 * ising_t *system: A struct
 *   enscapulating the information
 *   related to the system.
 *
 * returns
 * -------
 * float energy: The energy of the
 *   system in Joules.
 */
float energy_ising_t(ising_t *system)
{
  int length = system->length;
  int *ensemble = system->ensemble;
  float energy = 0.;

  for(int spin=0; spin<length; spin++)
  {
    energy -= ensemble[spin] *
      ensemble[modulo(spin+1,length)];
  }

  return energy;
}
```

You may notice that I am only counting the righthand neighbour of each spin. I chose this method because it is more optimal. If I was to count each neighbour then every pair would be counted twice and we would have to divide the final result by $2.$. By only counting one of the neighbours I have halved the number of computations.

A pair of spins is the *base unit* of the ising model so to calculate the entropy I counted the number of aligned pairs and then used the *chose* function to calculate the multiplicity. However, it was not quite this simples since $100!$ is $\sim 10^{157}$, which overflows an integer in the programs memory.

To fix this problem I used the stirling approximation to compute the entropy directly. This implies that the entropy should be less accurate at lower temperatures where the entropy is low and the Stirling approximation diverges from the actual entropy. The code that I used to calculate the entropy was,

```
/*
 * entropy_ising_t
 * ---------------
 * Calculate the entropy of a
 * configuration.
 *
 * parameters
 * ----------
 * ising_t *system: A struct
 *   enscapulating the information
 *   related to the system.
 *
 * returns
 * -------
 * float entropy: The entropy of the
 *   system in natural units.
```

```
*/
float entropy_ising_t(ising_t *system)
{
  int length = system->length;
  int *ensemble = system->ensemble;
  int up = 0;

  for (int spin=0; spin<length; spin++)
  {
    up += ensemble[spin] ==
      ensemble[modulo(spin+1,length)];
  }

  int down = length - up;

  float entropy = length * log(length) -
    up * log(up) - down * log(down);

  return entropy;
}
```


Figure 3: Histograms of the magnetisation for two ising systems of sizes $N = 100$ and $N = 500$. The temperatures are labelled in the titles and the y-axis is a unitless count.

Question 1 d)

Consider the energy depicted in figure 2. We see that the energy is a decreasing function of the temperature. This implies that the spins tend to align as the temperature decreases. This makes sense because the Boltzmann factor for the lower state becomes more favoured. Similarly the entropy is a increasing function of the temperature. As the spins tend to align at lower temperatures the number of ways to arrange the state becomes smaller, decreasing the entropy.

The heat capacity, also shown in figure 2 is interesting because it increases rapidly before asymptotically decreasing. If it was discontinuous then we would expect a phase transition however it is continuous. so we know that there is no phase transition. We know it is continuous because of equation 9.

Question 1 e)

Figure 3 showed me that the magnetisation was not zero as it was predicted to be at $0\epsilon/k$. I am not concerned because the distribution of samples, which roughly corresponds to a time average is symmetrically distributed around zero implying that the time average is identically zero. I would expect some variation of the results around zero as spins randomly flipped.

Since magnetisation sets in at a lower temperature for the larger sample I concluded that there was no phase transition. I decided this because the phase transition should occur at the same temperature for all lattice sizes. Taking the limit of a decreasing function presumably reduces to zero implying that there is no spontaneous magnetisation for the one-dimensional Ising model. This means that there is no phase transition.

Question 2 a)

For $\tau = 1.0\epsilon/k$ the random lattice very quickly relaxed into large *domains* of up and down spins. Over time the boundaries of the *domains* moved around the lattice coherently. I am/will using/use *coherent motion* to refer to a blob transporting across the network. For example, a group of down spins may start in the top right corner and remaining a single blob move to the middle of the system.

I noticed that given enough time the $\tau = 1.0\epsilon/k$ lattice relaxed until it was entirely one color. This is not shown in figure **??** simply because the lattice was not allowed to evolve for enough time. I used the same evolution time for the two dimensional Ising model that I used for the one dimensional Ising model, that is $1000N^2$ where I will use N to represent the side length of the system.

This is important because at lower temperatures the system took longer to equilibrate. As I saw for the one-dimensional case the *coherent domains* have a *meta-stable* lifetime before decaying. The shape of this domain influences the lifetime. The size seems to be a direct predictor of lifetime with the largest domain nearly always coming out on top. Another important feature that we noticed was that the coherent domains became particularly *meta-stable* if they formed rings.

The ring structures were *meta-stable* because of the cyclic boundary conditions. In a ring structure there are only two edges that are exposed to the other *meta-stable* state meaning that the anti-aligned interactions at the boundary are minimised. Additionally since the interaction depth of the Ising model is only nearest neighbour the *meta-stable* ring could exit in a band of just five spins. *The spins on the boundary have no idea how large their respective domains are.*

A final point to address for the $\tau = 1.0\epsilon/k$ model is the *time speckling*. As the system is allowed to evolve in equilibrium spins at random locations will randomly flip. This

is a byproduct of the metropolis algorithm but it is important because the *speckles* almost always remain a single spin. You will see later that this changes as the temperature is increased.

Heating the lattice to $\tau = 2.0\epsilon/k$ we see mush of the same behaviour. $\tau = 2.0\epsilon/k$ is still less than the critical temperature, so all of the same evolutionary events occur. *Meta-stable* domains form, move and die. Even the more dangerous *meta-stable* rings can form at this temperature. I would like to say that the lifetime of the *domains* has declined but I have no evidence.

The key difference between the $\tau 2.0\epsilon/k$ and the $\tau = 1.0\epsilon/k$ models is the *speckling*. At the higher temperatures the speckles occur more often and can result in very local short lived domains. This is why I postulated that the *domains* are shorter lived. The system is more active meaning that the spins on the boundary flip more readily. Moreover, the rapid speckling assists in the destruction of the *domain* when it occurs on the boundary.

Something worthy of note before we move above the condensation temperature is that *domains* die at there corners. This is because the corners represent an interface where a spin has more neighbours of one color than the other. This is another factor that prolongs the life of the *meta-stable* states, since they have no corners and minimised contact.

The behaviour was markedly different for the $\tau = 3.0\epsilon\ k$ model, because it is above the critical temperature. For starters the activity of the system has increased significantly. This means, among other things that the relaxation into equilibrium is almost instantaneous compared with the lower temperature models. In addition *domains* do not form at this temperature.

I have been using *domain* to refer to a structure that is *macroscopic* and *semi-permanent*. *Zonation* does occur at $\tau = 3.0\epsilon/k$ bur it does not meet these criteria. By this I mean that the zones are small ~ 10 spins maximum and temporary. Rather than considering the rapid zonation to be the evolution of the *domain* behaviour I consider it the evolution of the *speckling* behaviour.

At $\tau = 3.0\epsilon/k$ we also loose the *coherent* end state. At both of the other temperatures the final state was a completely magnetised block or a *meta-stable* ring. However, at $\tau = 3.0\epsilon/k$ there is no such state and the *zonation* appears to randomly move about the grid. At higher temperatures still the *zonation* is completely destroyed and not even microscopic domains are able to form.

Question 2 b)

Smoothly varying the temperature using I found that the *speckling* and *zonation* seems to dominate above $\sim 2.3\epsilon/k$ while below this *domains* start to form and given enough time the entire crystal becomes magnetised. The transition is smooth with respect to the *speckles*, which at very


| $T = 1.0$ | $T = 2.0$ | $T = 3.0$ |

Figure 4: Snapshots of the Ising system at multiple temperatures given a random initial starting position (top). You will notice that $\tau = 1.0\epsilon/k$ and $\tau = 2.0\epsilon/2$ have similar net magnetisation, but that $\tau = 2.0\epsilon/k$ has stronger *speckling* than $\tau = 1.0\epsilon/k$. You will notice that there is not net magnetisation for $\tau = 3.0\epsilon/k$ and that the *zonation* I described above is occurring.

low temperatures are very small and isolated with out a lifetime.

As I increased the temperature the *speckles* increases in size and in lifetime. A *speckle* is differentiated from a *zone* rather arbitrarily, the two main differences being transport and lifetime. The *zones* have a lifetime of long enough that they are able to *coherently move* across the crystal. On the other hand a *speckle* tends to die where it started and very quickly.

As we approach the critical temperature from below the *speckles* increase in size and disrupt the domains more quickly but, the system is still able to fully magnetise. At the critical temperature the *speckles* become so disruptive that the magnetisation cannot occur macroscopically but it still occurs locally in the form of the *zones* which outlive the now very active *speckles*.

Increasing the temperature further still the coherent *zones* disappear (visually) entirely and the speckles become the dominant behaviour. As a final note it is worth noting that these behaviours I have described were visible to me in part because of the settings of my simulation. I found it was best to run them with many steps per frame $5000 \sim 10000$ because it allows you to see the macroscopic behaviours.

I acknowledge that although I wanted to use the qt frame work to create my own gui I ran out of time. Instead I used the free online applet created by Daniel V. Schroeder for Weber state university to make the higher level observations. The code that I used to evolve the two dimensional system was,

/*

```
 * metropolis_step_ising_t
 * -----------------------
 * Evolve the system according to a
 * randomly weighted spin flip that
 * compares the probability of the
 * two states based on the Boltzmann
 * distribution of the two systems.
 *
 * parameters
 * ----------
 * ising_t *system: The system to
 *   evolve.
 */
void metropolis_step(ising_t *system)
{
  int length = system->length;
  int **ensemble = system->ensemble;
  float epsilon = system->epsilon;
  float temperature = system->temperature;
  float magnetic_field = system->magnetic_fi

  int row = random_index(length);
  int col = random_index(length);

  int spin = ensemble[row][col];
  int neighbours =
    ensemble[modulo(row + 1, length)][col] +
    ensemble[modulo(row - 1, length)][col] +
    ensemble[row][modulo(col + 1, length)] +
    ensemble[row][modulo(col - 1, length)];

  float magnetic_change = -2 * spin * magnet
  float interaction_change = 2 * epsilon * r
  float energy_change = magnetic_change + ir

  if ((energy_change < 0) ||
    (exp(- energy_change / temperature) > no
  {
    system -> ensemble[row][col] *= -1;
  }
}
```
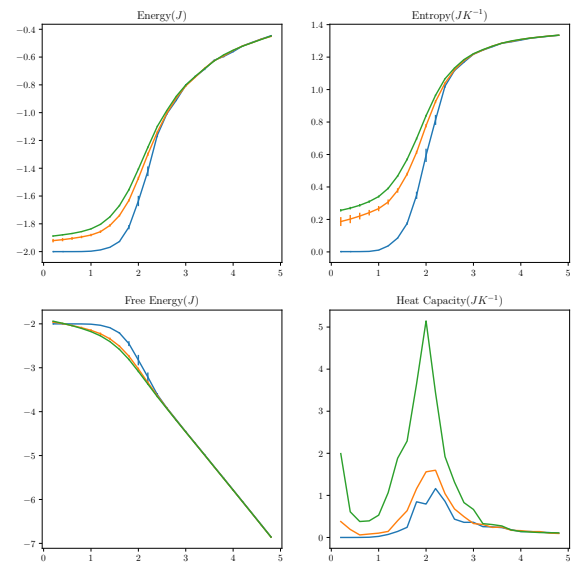
Figure 5:

Note: The `magnetic_field` was set to zero in the provided simulation and discussion, while the `epsilon` was set to one.

Question 2 c)